

Mostafa ElGhandour

Machine learning engineer Nano-degree

May 9<sup>th</sup>, 2019

## **Stock prediction using long-short term memory neural networks**

### **Project overview**

Predicting how the stock market will perform is one of the most difficult things to do. There are so many factors involved in the prediction – physical factors vs. physiological, rational and irrational behavior, etc. All these aspects combine to make share prices volatile and very difficult to predict with a high degree of accuracy, which makes it a very project to implement.

In this projected, I created an algorithm that can predict the price of the Egyptian pound relative to the US dollar.

### **Problem statement**

The goal was to create an algorithm that can learn and predict the pricing of the EGP-USD those were the tasks I followed for doing so:

1. Download and preprocess the EGP-USD pricing historical data.
2. Create a persistence model to use it as a benchmark.
3. Create a LSTM neural networks algorithm that predicts pricing.
4. Compare and tune the parameters for the LSTM to try to improve it more than the persistence model.

## Metrics

Mean squared error was used as an accuracy metric for both of the persistence model and the LSTM model

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2.$$

I created a persistence model to be used as a benchmark model by doing the following :

- 1.Transform the univariate dataset into a supervised learning problem.
- 2.Establish the train and test datasets for the test harness
- 3.Define the persistence model.
- 4.Make a forecast and establish a baseline performance
- 5.Plot the output

After creating the persistence model I calculated its MSE and used it as a reference for my model.

## Analysis

### Data exploration

The data I used was downloaded from a stock market website and I downloaded the data from 2000 but I used only the data after the floating (December 2016) of the Egyptian currency to remove any bias.

The data had the following fields:

- Date: it's the date of the current USD-EGP price.
- Price: the price of the conversion.


- Open: the opening price of the stock.
- High: the highest price the stock reached that day.
- Low: the lowest price the stock reached that day.
- Change: the change between high and low.

I disregarded the open, high, low, and change features and used the price feature only to simplify the problem, other features can be used in a later version.

### Exploratory visualization

The plot below shows the price of the USD-EGP starting from December 2016 to our current day, where the Y-axis represents the price of the Egyptian pound and the X-axis represents the time(day).

The days are simply translated into a sequenced number for easier understanding of the data, graphs and learning of the algorithm.



It can be seen that:

- The price doesn't change instantaneously (it takes a few days to spike) which is very good in our case in using the short memory technique.

## Algorithms and techniques

The algorithm I used is a LSTM neural network which is a powerful type of neural network designed to handle sequence dependence is called recurrent neural networks

The Long Short-Term Memory network or LSTM network is a type of recurrent neural network used in deep learning because very large architectures can be successfully trained.

The following parameters were tuned to optimize the classifier:

- Number of layers for the neural network
- Layer types
- Layer parameters
- Number of epochs, batch size.

To tune the epochs and batch size, I used grid search to optimize the parameters which in return gave me the best parameters combination from the 2 parameter arrays I entered.

## Benchmark

Establishing a baseline is essential on any time series forecasting problem, The most common baseline method for supervised machine learning is the Rule algorithm.

This algorithm predicts the majority class in the case of classification, or the average outcome in the case of regression. This could be used for time series, but does not respect the serial correlation structure in time series datasets. The equivalent technique for use with time series dataset is the persistence algorithm.

So, I created my own persistence model to use it as a benchmark,

The steps:

- Create a lagged representation of the dataset.
- Split the dataset into train and testing data.
- Train the data using the original and the lagged dataset.
- Calculate the MSE of the model and use it as a reference.

Having done these steps, I got an MSE of 0.002 which was used as a reference for my model.

## Methodology

### **Data preprocessing**

The steps were as follows:

1. The dataset was read using pandas.
2. The price column was taken into a “working data” list.
3. All data before December 2016(floating of the Egyptian pound) was removed.
4. The data was split into training and test sets with respect to time (randomizing the training and testing set would be very hard to predict in a time series model).

The data preprocessing was pretty simple and straight forward as my data is not complicated at all.

## Implementation

The implementation was as follows:

1. Preprocessing the data:
  - a. the dataset was read using pandas.
  - b. The price “column” was put into a “working\_data” list.
  - c. The dataset was plotted so it could be further analyzed.
2. Creating the persistence model:
  - a. A lagged representation of the “price” column was created using (concat).
  - b. Splitting the data into train and test sets.
  - c. The mean squared error of the data and its lagged representation was calculated
3. Creating the LSTM model:
  - a. Split the data into train and test sets
  - b. Create a function that creates another dataset in a lagged representation, 1-time step was chosen to be able to compare it to the persistence model.
  - c. Create a neural network with 1 input layer and a hidden layer with 4 neurons and an output layer, I tried to further optimize the layers, but this was the best results I got.
  - d. Create a grid search to find the best parameters for the following neural network
  - e. Use the optimized parameters (batch size and epochs) into my neural network.

- f. Compare and plot the result of the trained model.


## Refinement

Grid search was used to tune the parameters, as tuning them manually was almost impossible so I created a set of number of the epochs (10,25,50,100, 500, 250) and a set of parameters of the batch size(1,10, 20, 40, 60, 80, 100) and using the gridsearch I got the optimal parameters which were 25 epoch and 1 batch size after 3 hours of processing.

## Results

### Model evaluation, validation and justification

The final tested model gave a test score on the mean squared error of 0.05 which is more than what I got in the persistence model, because I could not improve it more than that.




The blue line represents the original data, the orange is the predicted training set and the green is the predicted test set,

Shown in the figure above how accurately my model can predict the training and testing data.

The model can be trusted and used as a stock predication mode with a little bit of error that I think can be adjusted with a little bit of parameter tuning and a stronger pc for gridsearch.

## Conclusion

### Free-Form visualization



as shown in the above figure, the model does an excellent job of fitting the train and test datasets but still with a larger error than the persistence model, it needs some parameter tuning to improve it.

### Reflection

To summarize the process done in this project we can simply look at these steps :

1. An initial problem and relevant dataset were chosen.
2. The dataset was downloaded and preprocessed.



3. A persistence model was created to be used as a benchmark model and its mean squared error was calculated.
4. A LSTM neural network was created.
5. A gridsearch algorithm was implemented to tune the parameters for the neural network.
6. The LSTM was tuned using the output results from the gridsearch.
7. The mean squared error from the trained LSTM model was calculated and compared to the persistence model's.

Creation and tuning of the LSTM was the hardest part as I could not put my hands on how to split the data into labels and features, and the networks layers were really hard to play with as they would just worsen the results whenever I tried adding/removing layers.

## **Improvement**

of course there's a large room for improvement in this algorithm so it can be used as a legitimate algorithm for predicting stock prices, and they can be summarized in the following:

1. better hardware to ease the process of using gridsearch and using more sets of parameters.
2. More features can be used in the model to improve its accuracy.
3. More lagged time steps should be used rather the 1-time step only I used, that could help the algorithm be more accurate in predicting the far future.
4. And of course an easy GUI could be used for prediction.

## References

1. “Time Series Prediction with LSTM Recurrent Neural Networks in Python with Keras.” *Machine Learning Mastery*, 26 Apr. 2019, [machinelearningmastery.com/time-series-prediction-lstm-recurrent-neural-networks-python-keras/](https://machinelearningmastery.com/time-series-prediction-lstm-recurrent-neural-networks-python-keras/).
2. “How to Grid Search Hyperparameters for Deep Learning Models in Python With Keras.” *Machine Learning Mastery*, 9 May 2019, [machinelearningmastery.com/grid-search-hyperparameters-deep-learning-models-python-keras/](https://machinelearningmastery.com/grid-search-hyperparameters-deep-learning-models-python-keras/).
3. Nayak, Asutosh. “Predicting Stock Price with LSTM.” *Towards Data Science*, Towards Data Science, 18 Mar. 2019, [towardsdatascience.com/predicting-stock-price-with-lstm-13af86a74944](https://towardsdatascience.com/predicting-stock-price-with-lstm-13af86a74944).
4. Nayak, Asutosh. “Finding the Right Architecture for Neural Network.” *Towards Data Science*, Towards Data Science, 18 Mar. 2019, [towardsdatascience.com/finding-the-right-architecture-for-neural-network-b0439efa4587](https://towardsdatascience.com/finding-the-right-architecture-for-neural-network-b0439efa4587).
5. “How to Make Baseline Predictions for Time Series Forecasting with Python.” *Machine Learning Mastery*, 26 Apr. 2019, [machinelearningmastery.com/persistence-time-series-forecasting-with-python/](https://machinelearningmastery.com/persistence-time-series-forecasting-with-python/).